

18.335 Take-Home Midterm Exam Solutions: Spring 2021

Problem 1: (34 points)

As suggested, let's define a recurrence for Horner's rule so that $s_k = c_k + xs_{k+1}$ with $s_{n-1} = c_{n-1}$, so that $f = s_0$, and hence the corresponding floating-point algorithm for inputs in \mathbb{F} (as in the summation notes from class, this is just a convenience — inputs in \mathbb{R} would just give some additional $1 + \epsilon$ factors that we could incorporate into the same analysis) is

$$\tilde{s}_{n-1} = c_{n-1}$$

$$\tilde{s}_k = c_k \oplus (x \otimes \tilde{s}_{k+1}) = [c_k + x\tilde{s}_{k+1}(1 + \epsilon_k)](1 + \epsilon'_k)$$

for some $|\epsilon_k| \leq \epsilon_{\text{machine}}$ and $|\epsilon'_k| \leq \epsilon_{\text{machine}}$, by the fundamental axiom of floating-point arithmetic. Hence

$$\tilde{s}_k = \underbrace{(c_k + c_k \epsilon'_k + x\tilde{s}_{k+1}[\epsilon_k + \epsilon'_k])}_{\tilde{c}_k} + x\tilde{s}_{k+1} = \sum_{j=k}^{n-1} \tilde{c}_j x^{j-k},$$

where we have defined \tilde{c}_k . By construction

$$\tilde{f}(c, x) = \tilde{s}_0 = f(\tilde{c}, x),$$

i.e. the output of \tilde{f} is equal to the exact polynomial output with altered coefficients \tilde{c} . To establish backwards stability, we now need to show that $\|\tilde{c} - c\| = \|c\|O(\epsilon_{\text{machine}})$ in some norm. Similar to the summation analysis in class, we will choose the L^1 norm for convenience. In particular,

$$|\tilde{c}_k - c_k| \leq |c_k| |\epsilon'_k| + |\epsilon_k + \epsilon'_k| |x| |\tilde{s}_{k+1}| + O(\epsilon_{\text{machine}}^2).$$

Now, use the fact that

$$|\tilde{s}_k| = \left| \sum_{j=k}^{n-1} \tilde{c}_j x^{j-k} \right| \leq \left| \sum_{j=k}^{n-1} \tilde{c}_j \right| \left(\max\{1, |x|^{n-k}\} \right) \leq \|\tilde{c}\|_1 \left(\max\{1, |x|^{n-1}\} \right)$$

to obtain

$$|\tilde{c}_k - c_k| \leq |c_k| |\epsilon'_k| + |\epsilon_k + \epsilon'_k| |x| \|\tilde{c}\|_1 \left(\max\{1, |x|^{n-1}\} \right) + O(\epsilon_{\text{machine}}^2) \leq |c_k| \epsilon_{\text{machine}} + 2|x| \epsilon_{\text{machine}} \|\tilde{c}\|_1 \left(\max\{1, |x|^{n-1}\} \right) + O(\epsilon_{\text{machine}}^2)$$

and hence

$$\|\tilde{c} - c\|_1 \leq \|c\|_1 \epsilon_{\text{machine}} + 2n|x| \epsilon_{\text{machine}} \|\tilde{c}\|_1 \left(\max\{1, |x|^{n-1}\} \right) + O(\epsilon_{\text{machine}}^2) = \|c\|_1 O(\epsilon_{\text{machine}}) + \|\tilde{c}\|_1 O(\epsilon_{\text{machine}})$$

where we can put any higher-order terms and c -independent coefficients into the $\epsilon_{\text{machine}}$. Finally, exactly as in Sec. 5.1 of the summation-stability notes from class, $\|\tilde{c}\|_1 O(\epsilon_{\text{machine}}) = \|c\|_1 O(\epsilon_{\text{machine}})$, because the difference is higher-order in $\epsilon_{\text{machine}}$, so we finally have

$$\|\tilde{c} - c\|_1 = \|c\|_1 O(\epsilon_{\text{machine}})$$

and our backwards-stability proof is complete.

Problem 2: (20+13 points)

Suppose that

$$T = \begin{pmatrix} \alpha_1 & \beta_1 & & & \\ \overline{\beta_1} & \alpha_2 & \beta_2 & & \\ & \overline{\beta_2} & \ddots & \ddots & \\ & & \ddots & \alpha_{m-1} & \beta_{m-1} \\ & & & \overline{\beta_{m-1}} & \alpha_m \end{pmatrix}$$

is an $m \times m$ Hermitian ($T = T^*$) complex tridiagonal matrix.

(a) Let

$$D = \begin{pmatrix} d_1 & & & \\ & d_2 & & \\ & & \ddots & \\ & & & d_m \end{pmatrix}.$$

Then

$$\hat{T} = D^{-1}TD = \begin{pmatrix} \alpha_1 & \beta_1 \frac{d_2}{d_1} & & \\ \overline{\beta_1} \frac{d_1}{d_2} & \alpha_2 & \beta_2 \frac{d_3}{d_2} & \\ & \overline{\beta_2} \frac{d_2}{d_3} & \ddots & \\ & & \ddots & \alpha_{m-1} & \beta_{m-1} \frac{d_m}{d_{m-1}} \\ & & & \overline{\beta_{m-1}} \frac{d_{m-1}}{d_m} & \alpha_m \end{pmatrix}.$$

Note that α_k is already real since $T = T^*$, so we only need to choose D to make the off-diagonal terms real. Write β_k in polar form as $\beta_k = r_k e^{i\phi_k}$. In $D^{-1}TD$, this is rescaled in the upper diagonal to $\beta_k \frac{d_{k+1}}{d_k}$.

To make this purely real, we therefore want $\frac{d_{k+1}}{d_k} = e^{-i\phi_k}$ (or $-e^{-i\phi_k}$ would also work). Hence, let us define the diagonals by the recurrence:

$$\begin{aligned} d_1 &= 1 \text{ (arbitrary choice),} \\ d_{k+1} &= d_k e^{-i\phi_k} = e^{-i\sum_{j \leq k} \phi_j}. \end{aligned}$$

This also fixes the lower diagonal, since $\frac{d_k}{d_{k+1}} = e^{+i\phi_k}$ and hence:

$$\overline{\beta_k} \frac{d_k}{d_{k+1}} = r_k e^{-i\phi_k} \frac{d_k}{d_{k+1}} = r_k.$$

D is clearly unitary since the diagonal entries all have $|d_k| = 1$ (as long as we chose $|d_1| = 1$).

- (b) We should apply QR iterations to \hat{T} , which gives the same result because T and \hat{T} are **similar** matrices (differing only by a unitary change of basis: they have the same eigenvalues, and the eigenvectors differ by a factor of D), and operating on \hat{T} is faster because then the QR factors will also be purely real, and operations on real numbers take fewer floating-point operations than corresponding operations on complex numbers (adding complex numbers takes 2 flops and multiplying them takes 6 flops, versus 1 addition and 1 multiplication, respectively, for real numbers). (It also requires half as much memory and hence will be more cache-efficient.)

Problem 3: (19+14 points)

Suppose A is an $m \times n$ matrix with $n > m$ and rank m (linearly independent rows): a “wide” matrix. In this case, $Ax = b$ has infinitely many solutions x for any right-hand side $b \in \mathbb{C}^m$ (it is an *underdetermined* system of equations). We compute the QR factorization of the conjugate-transpose $A^* = QR = \hat{Q}\hat{R}$ by some backwards-stable algorithm (e.g. Householder QR), where Q is $n \times n$ and R is $n \times m$, and \hat{Q} is the “thin” QR (the first m columns of Q) and \hat{R} is correspondingly the first m rows of R .

- (a) As suggested by the hint, let’s write the solutions in the Q basis as

$$x = Qy = \begin{pmatrix} \hat{Q} & Q_{\perp} \end{pmatrix} \begin{pmatrix} \hat{y} \\ y_{\perp} \end{pmatrix} = \hat{Q}\hat{y} + Q_{\perp}y_{\perp},$$

where we have broken up y into the m coefficients of \hat{Q} and the $n - m$ coefficients of Q_{\perp} . Then

$$b = Ax = \hat{R}^* \hat{Q}^* x = \hat{R}^* \hat{Q}^* (\hat{Q}\hat{y} + Q_{\perp}y_{\perp}) = \hat{R}^* \hat{y},$$

since $\hat{Q}^* \hat{Q} = I$ (being an orthonormal basis) and $\hat{Q}^* Q_\perp = 0$ since Q_\perp spans the orthogonal complement of \hat{Q} , which spans the row space $C(A^*)$, and hence Q_\perp spans $C(A^*)^\perp = N(A)$. This tells us several things.

- (i) \hat{y} is uniquely determined by solving the $m \times m$ lower-triangular system $\hat{R}^* \hat{y} = b$.
 - (ii) any value of y_\perp is allowed. (Equivalently, the term $Q_\perp y_\perp$ represents anything in the null space of A .)
 - (iii) $\|x\|_2^2 = (\hat{Q}\hat{y} + Q_\perp y_\perp)^* (\hat{Q}\hat{y} + Q_\perp y_\perp) = \|\hat{y}\|_2^2 + \|y_\perp\|_2^2 \geq \|\hat{y}\|_2^2$, which is clearly minimized for $y_\perp = 0$.
 - (iv) Hence the minimum-norm solution is $x = \hat{Q}\hat{y} = \hat{Q}(\hat{R}^*)^{-1}b$.
- (b) As suggested, suppose we are *given* the $A^* = QR$ factorization computed by Householder QR, and now we want to work out the additional cost of computing our minimum-norm solution x above. Then the additional cost is to (i) solve $\hat{R}^* \hat{y} = b$ and (ii) multiply $\hat{Q}\hat{y}$. Computation (i) can use forward-substitution since \hat{R}^* is lower-triangular, so it is $\Theta(m^2)$.

For computation (ii), remember that Householder QR does not actually compute Q explicitly — rather, it computes a set of reflectors and provides an efficient method (algorithm 10.3 in the book) to multiply Q by vectors as a sequence of reflections. Here, we need to apply that algorithm 10.3, but we are multiplying by $y = \begin{pmatrix} \hat{y} \\ 0 \end{pmatrix}$ so we can maybe save a few operations by skipping multiplications with the zero components. Algorithm 10.3 initializes $x = y$ and then computes $x_{k:n} = x_{k:n} - 2v_k(v_k^* x_{k:n})$ for $k = m$ down to 1 (where I've swapped m and n compared to the book). Here, we hope that this simplifies because $y_{m+1:n} = 0$. However, after the *very first* ($k = m$) step of the algorithm, in which $v_m^* x_{m:n} = v_m[m] \hat{y}[m] \neq 0$ in general, we subtract a multiple of $2v_k$ from $x_{m:n}$ and *all* the components become nonzero for subsequent steps. So, in the end the savings from the zero entries of y are asymptotically negligible and the cost of step (ii) is $\Theta(nm)$ exactly as in the book.

(Note that even if you assume we have \hat{Q} explicitly, e.g. if you used modified Gram–Schmidt, multiplying it by a vector still has $\Theta(nm)$ cost.)

Hence, the overall complexity is $\Theta(m^2) + \Theta(nm) = \Theta(mn)$: since $n > m$ the step (ii) dominates.